

Problem Statement: Abalone Age Prediction

Description:- Predicting the age of abalone from physical measurements. The age of abalone is determined by cutting the shell through the cone, staining it, and counting the number of rings through a microscope -- a boring and time-consuming task. Other measurements, which are easier to obtain, are used to predict age. Further information, such as weather patterns and location (hence food availability) may be required to solve the problem.

Building a Regression Model

1. Download the dataset
2. Load the dataset into the tool.
3. Perform Below Visualizations.
 - Univariate Analysis
 - Bi-Variate Analysis
 - Multi-Variate Analysis
4. Perform descriptive statistics on the dataset.
5. Check for Missing values and deal with them.
6. Find the outliers and replace them outliers
7. Check for Categorical columns and perform encoding.
8. Split the data into dependent and independent variables.
9. Scale the independent variables
10. Split the data into training and testing

- ```
In []: #import libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb
import plotly.express as px
```

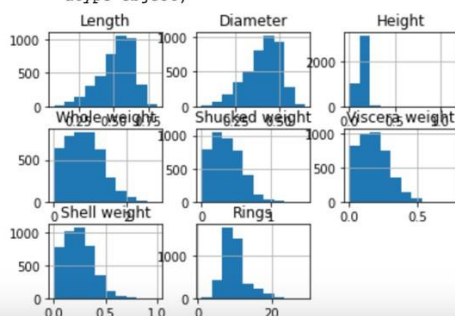
```
In []: data = pd.read_csv('/content/drive/My Drive/Machine Learning/abalone.csv')
 data
```

4177 rows x 9 columns

- Univariate Analysis

```
In []: data['Rings'].value_counts()
data.hist()
```


```
Out[]: array([[,
 ,
],
 [,
 ,
],
 [,
 ,
]],
 dtype=object)
```



- Bi-Variate Analysis

```
In []: plt.scatter(data.Rings, data.Sex)
plt.title('The Gender of Abalone vs Number of Rings')
plt.xlabel('No. of Rings')
plt.ylabel('Gender')
```


```
Out[]: Text(0, 0.5, 'Gender')
```



- Multi-Variate Analysis

```
In []: sb.heatmap(data.corr(),annot=True)
```

```
Out[]:
```



## 4. Perform descriptive statistics on the dataset.

```
In []: data.info()
```

```
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
Column Non-Null Count Dtype
--- -
0 Sex 4177 non-null object
1 Length 4177 non-null float64
2 Diameter 4177 non-null float64
3 Height 4177 non-null float64
4 Whole weight 4177 non-null float64
5 Shucked weight 4177 non-null float64
6 Viscera weight 4177 non-null float64
7 Shell weight 4177 non-null float64
8 Rings 4177 non-null int64
dtypes: float64(7), int64(1), object(1)
memory usage: 293.8+ KB
```

```
In []: data.describe()
```

```
Out[]:
```

|       | Length      | Diameter    | Height      | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings       |
|-------|-------------|-------------|-------------|--------------|----------------|----------------|--------------|-------------|
| count | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000  | 4177.000000    | 4177.000000    | 4177.000000  | 4177.000000 |
| mean  | 0.523992    | 0.407881    | 0.139516    | 0.828742     | 0.359367       | 0.180594       | 0.238831     | 9.933684    |
| std   | 0.120093    | 0.099240    | 0.041827    | 0.490389     | 0.221963       | 0.109614       | 0.139203     | 3.224169    |
| min   | 0.075000    | 0.055000    | 0.000000    | 0.002000     | 0.001000       | 0.000500       | 0.001500     | 1.000000    |
| 25%   | 0.450000    | 0.350000    | 0.115000    | 0.441500     | 0.186000       | 0.093500       | 0.130000     | 8.000000    |
| 50%   | 0.545000    | 0.425000    | 0.140000    | 0.799500     | 0.336000       | 0.171000       | 0.234000     | 9.000000    |
| 75%   | 0.615000    | 0.480000    | 0.165000    | 1.153000     | 0.502000       | 0.253000       | 0.329000     | 11.000000   |
| max   | 0.815000    | 0.650000    | 1.130000    | 2.825500     | 1.488000       | 0.760000       | 1.005000     | 29.000000   |

## 5. Check for Missing values and deal with them.

There is no missing values

```
In []: data.isnull().any()
```

```
Out[]: Sex False
Length False
Diameter False
Height False
Whole weight False
Shucked weight False
Viscera weight False
Shell weight False
Rings False
dtype: bool
```

## 6. Find the outliers and replace them outliers

The dataset does not have a outliers

```
In []: fig = px.histogram(data, x='Whole weight')
fig.show()
```

## 7. Check for Categorical columns and perform encoding.

There is one Categorical column SEX is replaced by an Integer

```
In []: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
data["Sex"] = le.fit_transform(data["Sex"])
data["Sex"]
```

```
Out[]: 0 2
1 2
2 0
3 2
4 1
..
4172 0
4173 2
4174 2
4175 0
4176 2
Name: Sex, Length: 4177, dtype: int64
```

## 8. Split the data into dependent and independent variables.

```
In []: x=data.iloc[:,0:8].values
 y=data.iloc[:,8:9].values
```

```
In []: x
```

```
Out[]: array([[2. , 0.455 , 0.365 , ..., 0.2245, 0.101 , 0.15],
 [2. , 0.35 , 0.265 , ..., 0.0995, 0.0485, 0.07],
 [0. , 0.53 , 0.42 , ..., 0.2565, 0.1415, 0.21],
 ...,
 [2. , 0.6 , 0.475 , ..., 0.5255, 0.2875, 0.308],
 [0. , 0.625 , 0.485 , ..., 0.531 , 0.261 , 0.296],
 [2. , 0.71 , 0.555 , ..., 0.9455, 0.3765, 0.495]])
```

```
In []: y
```

```
Out[]: array([[15],
 [7],
 [9],
 ...,
 [9],
 [10],
 [12]])
```

## 9. Scale the independent variables

```
In []: x=data.iloc[:,0:8]
 print(x.head())
```

|   | Sex | Length | Diameter | Height | Whole weight | Shucked weight | \ |
|---|-----|--------|----------|--------|--------------|----------------|---|
| 0 | 2   | 0.455  | 0.365    | 0.095  | 0.5140       | 0.2245         |   |
| 1 | 2   | 0.350  | 0.265    | 0.090  | 0.2255       | 0.0995         |   |
| 2 | 0   | 0.530  | 0.420    | 0.135  | 0.6770       | 0.2565         |   |
| 3 | 2   | 0.440  | 0.365    | 0.125  | 0.5160       | 0.2155         |   |
| 4 | 1   | 0.330  | 0.255    | 0.080  | 0.2050       | 0.0895         |   |

|   | Viscera weight | Shell weight |
|---|----------------|--------------|
| 0 | 0.1010         | 0.150        |
| 1 | 0.0485         | 0.070        |
| 2 | 0.1415         | 0.210        |
| 3 | 0.1140         | 0.155        |
| 4 | 0.0395         | 0.055        |

## 10. Split the data into training and testing

```
In []: from sklearn.model_selection import train_test_split
 x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3,random_state=0)
```

```
In []: x_train.shape
```

```
Out[]: (2923, 8)
```

```
In []: x_test.shape
```

```
Out[]: (836, 8)
```

Out[ ]: (836, 8)

## 11. Build the Model

```
In []: from sklearn.linear_model import LinearRegression
lr = LinearRegression()
```

## 12. Train the Model

```
In []: lr.fit(x_train, y_train)
```

Out[ ]: LinearRegression()

## 13. Test the Model

```
In []: y_pred = lr.predict(x_test)
print((y_test)[0:6])
print((y_pred)[0:6])
```

```
[[13]
 [8]
 [11]
 [5]
 [12]
 [11]]
[[13.11640829]
 [9.65691091]
 [10.35350972]
 [5.63648715]
 [10.67436485]
 [11.95341338]]
```

## 14. Measure the performance using Metrics.

```
In []: # RMSE(Root Mean Square Error)

from sklearn.metrics import mean_squared_error
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
print("RMSE value : {:.2f}".format(rmse))
```

RMSE value : 2.26

```
In []: from sklearn.model_selection import cross_val_score
cv_scores = cross_val_score(lr, x, y, cv=5)
sco=cv_scores.round(4)
print(cv_scores.round(4))
print("Average", sco.sum()/5)
```

```
[0.4113 0.1574 0.4807 0.5046 0.4362]
Average 0.39803999999999995
```